

Lock Service *(simple Chubby)*

*Distributed Lock Service ; NuRaft ; In-memory log;
Persistent State Machine;*

*Aanandita Dhawan
dhawan6@wisc.edu*

*Evan Wireman
wireman@wisc.edu*

*Safi Nassar
nassar2@wisc.edu*

Design Decisions

- NuRaft for consistency
 - Open source, simple to implement, easy to import, leader consensus done behind the scenes
- No subscription to events
 - Would have required complicated heartbeat RPC responses
- `msd.channels.h` was used for inter-thread communication
 - Blocking channel implementation, open source
- Single binary, separate namespaces for server and client
 - A Client can include `app.h` and have access to client methods

NuRaft

- NuRaft by ebay (open source) for consensus
- Easy to configure and install
- Simple API (get, set)
- Shrouds leader election
- Did not require much code alteration to be able to log [file_name, contents] pairs.
- Asynchronous mode (we used synchronous for simplicity)

NuRaft Flow

User	Leader	Follower(s)
X----->		raft_server::append_entries()
	X	log_store::append()
	X	state_machine::pre_commit()
<-----	(X)	(async_handler mode) return raft_server::append_entries()
	X----->	Send logs
		(X) (if conflict) state_machine::rollback()
		(X) (if conflict) log_store::write_at()
		(X) (if conflict) state_machine::pre_commit()
		X log_store::append()
		X state_machine::pre_commit()
		(X) (commit of previous logs) state_machine::commit()
	<-----	X Respond
	X	RESULT <- state_machine::commit()
<-----	(X)	(blocking mode) return raft_server::append_entries()
		with RESULT
		(async_handler mode) invoke user-defined handler
		with RESULT

Structs

```
struct Lock {
    std::string path;                // Path to the file
    LockStatus status;              // Is it shared or exclusive?
    std::shared_ptr<std::map<std::string, bool>> owners; // Who owns the lock
    std::string content;            // File content
};

struct Session {
    string client_id;               // Id of the client with whom this session exists
    chrono::system_clock::time_point start_time; // Start time of the local session
    chrono::milliseconds lease_length; // Length of the session lease
    shared_ptr<msd::channel<int>> block_reply; // Channel used for blocking the reply to keep_alive rpcs
    shared_ptr<map<string, shared_ptr<Lock>>> locks; // Locks acquired with the session
    bool terminated;                // Indicator of if the session has been terminated manually
}

enum LockStatus {
    EXCLUSIVE,
    SHARED,
    FREE
}
```

Data structures *(singleton pattern)*

Server:

```
map<file_path, struct Lock> locks;  
map<client_id, struct Session> sessions;
```



Client

```
string session_id;  
chrono::system_clock::time_point lease_start;  
chrono::milliseconds lease_length;  
shared_ptr<map<string, LockStatus>> locks;  
bool jeopardy;  
bool expired;  
shared_ptr<Node> master; // A Node is just an abstraction for rpc endpoint
```

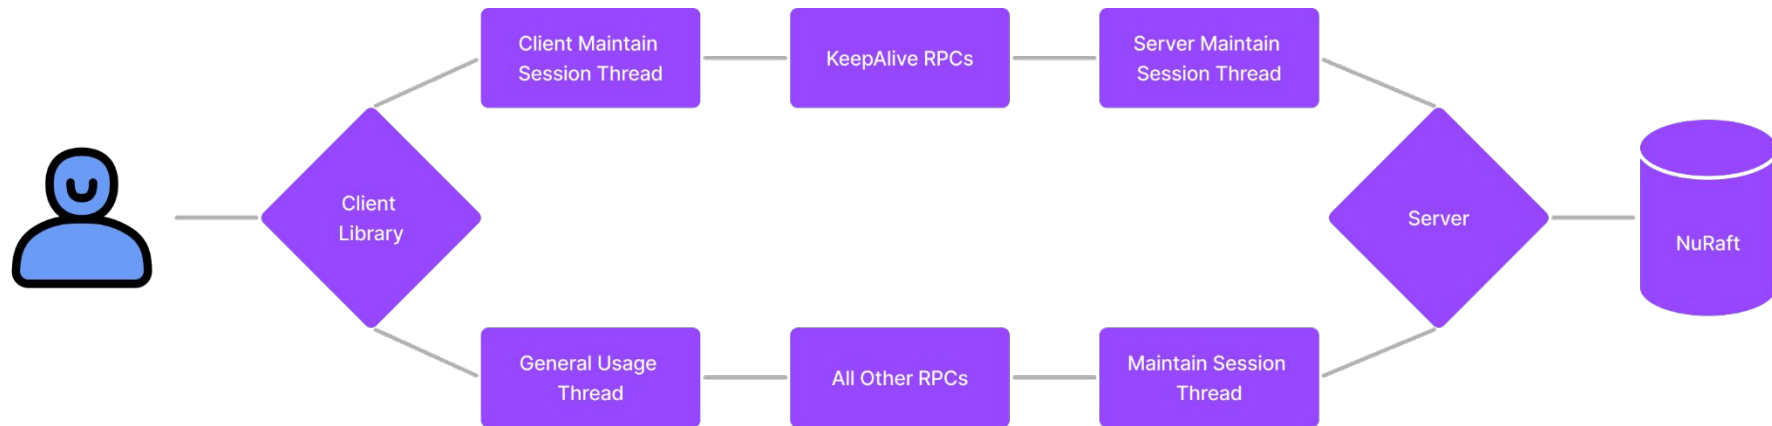
RPCs

- `Init_session(client_id)`
- `Close_session(client_id)`
- `Keep_alive(client_id, time_point)`
- `Keep_alive(client_id, time_point, map<file_path, LockStatus>)`
- `Open_lock(client_id, file_path)`
- `Delete_lock(client_id, file_path)`
- `Acquire_lock(client_id, file_path, LockMode)`
- `Release_lock(client_id, file_path)`
- `Read(client_id, file_path)`
- `Write(client_id, file_path, content)`

Chubby Architecture



Our Project's Architecture



Live Demo

We will show our implementation, in a 5 node cluster:

- Create and maintain sessions (heartbeats)
- Handle server failures (client in jeopardy)
- Handle client failures (session timeout)
- Open/close locks
- Acquire/release locks
- Read/write to locks

Takeaways

Future work:

- Event subscription
- Evaluation using raft vs. paxos for consensus
- Comparison with ZooKeeper

Lessons Learned:

- The architecture of Chubby in great detail
- Tailoring open source projects (NuRaft, msd channels) to match our needs
- Test-first development can be useful for projects that has a lot of room for bugs